

# Data Mining Project

## Data Description

This dataset contains comprehensive information on 2,392 high school students, detailing their demographics, study habits, parental involvement, extracurricular activities, and academic performance. The target variable, GradeClass, classifies students' grades into distinct categories.

### Variables:

- **StudentID:** A unique identifier assigned to each student (1001 to 3392).
- **Age:** The age of the students ranges from 15 to 18 years.
- **Gender:** Gender of the students, where 0 represents Male and 1 represents Female.
- **Ethnicity:** The ethnicity of the students, coded as follows:
  - 0: Caucasian
  - 1: African American
  - 2: Asian
  - 3: Other
- **ParentalEducation:** The education level of the parents, coded as follows:
  - 0: None
  - 1: High School
  - 2: Some College
  - 3: Bachelor's
  - 4: Higher
- **StudyTimeWeekly:** Weekly study time in hours, ranging from 0 to 20.
- **Absences:** Number of absences during the school year, ranging from 0 to 30.
- **Tutoring:** Tutoring status, where 0 indicates No and 1 indicates Yes.

- **ParentalSupport:** The level of parental support, coded as follows:
  - 0: None
  - 1: Low
  - 2: Moderate
  - 3: High
  - 4: Very High
- **Extracurricular:** Participation in extracurricular activities, where 0 indicates No and 1 indicates Yes.
- **Sports:** Participation in sports, where 0 indicates No and 1 indicates Yes.
- **Music:** Participation in music activities, where 0 indicates No and 1 indicates Yes.
- **Volunteering:** Participation in volunteering, where 0 indicates No and 1 indicates Yes.
- **GPA:** Grade Point Average on a scale from 2.0 to 4.0, influenced by study habits, parental involvement, and extracurricular activities.
- **GradeClass:** Classification of students' grades based on GPA:
  - 0: 'A' (GPA  $\geq 3.5$ )
  - 1: 'B' ( $3.0 \leq \text{GPA} < 3.5$ )
  - 2: 'C' ( $2.5 \leq \text{GPA} < 3.0$ )
  - 3: 'D' ( $2.0 \leq \text{GPA} < 2.5$ )
  - 4: 'F' (GPA  $< 2.0$ )

## Data Importing:

```
mydata0=read.csv("C:\\Users\\Lenovo\\Downloads\\archive (1)\\Student_performa
nce_data _.csv")
#View(mydata0)
mydata0$GradeClass=ifelse(mydata0$GradeClass==0,"A",
                           ifelse(mydata0$GradeClass==1,"B",
                                   ifelse(mydata0$GradeClass==2,"C",
                                           ifelse(mydata0$GradeClass==3,"D",
                                                   ifelse(mydata0$GradeClass==4,"F",NA))))))
mydata=mydata0[,-c(1,14)]
dim(mydata)
```

```
## [1] 2392 13
```

```
Y=mydata[, "GradeClass"]
```

```
X=mydata[, -13]
```


## Exploratory Data Analysis

```
sum(is.na(mydata))
```



```
## [1] 0
```



We confirm from the info that there are no nulls in this DataFrame.

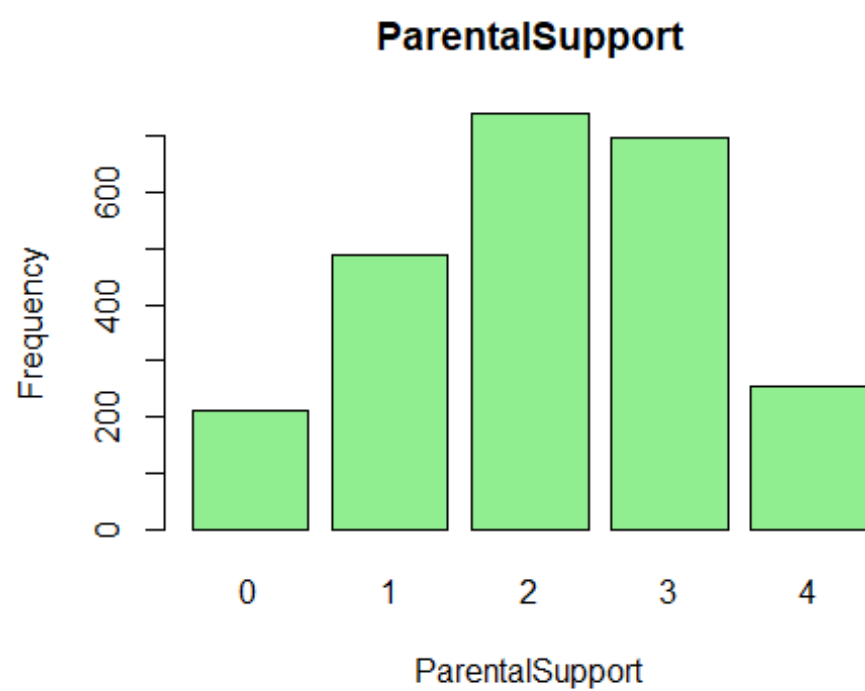
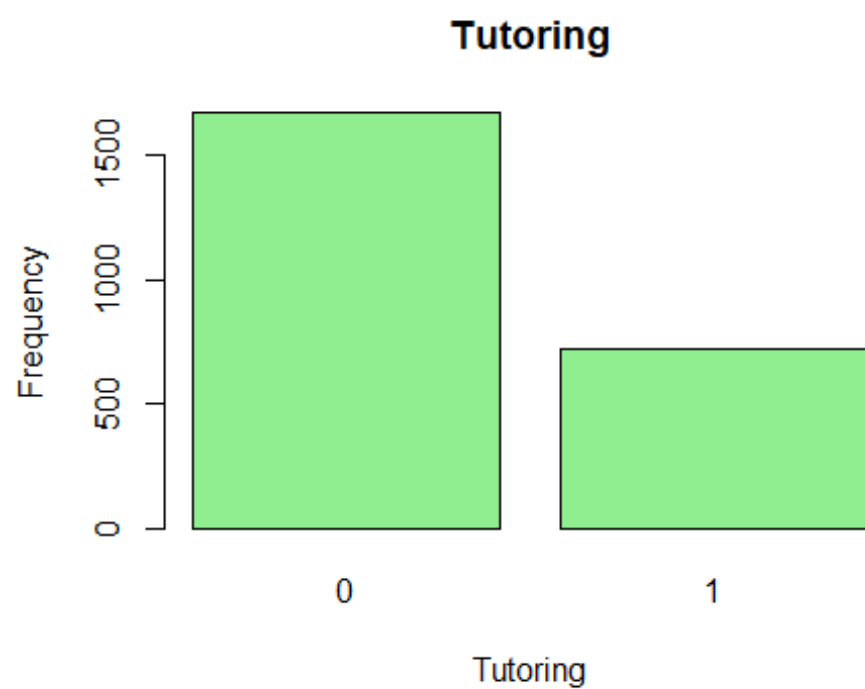
```
a=table(mydata$GradeClass)
label=c("A", "B", "C", "D", "F")
pie(a, labels=paste0(label, "(", round(a/ sum(a)*100, 1), "%)"),
    main = "Pie Chart of GradeClass")
```

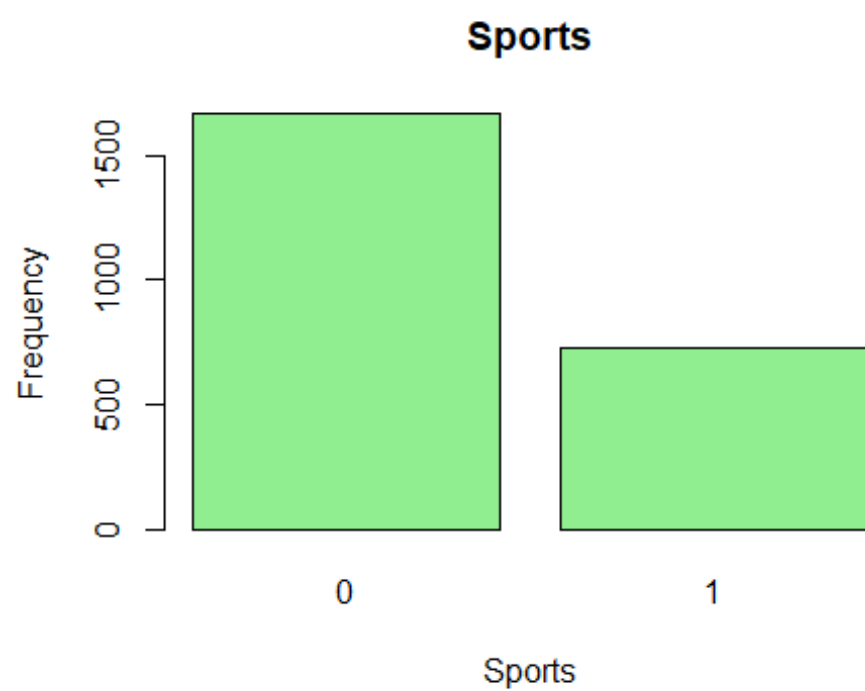
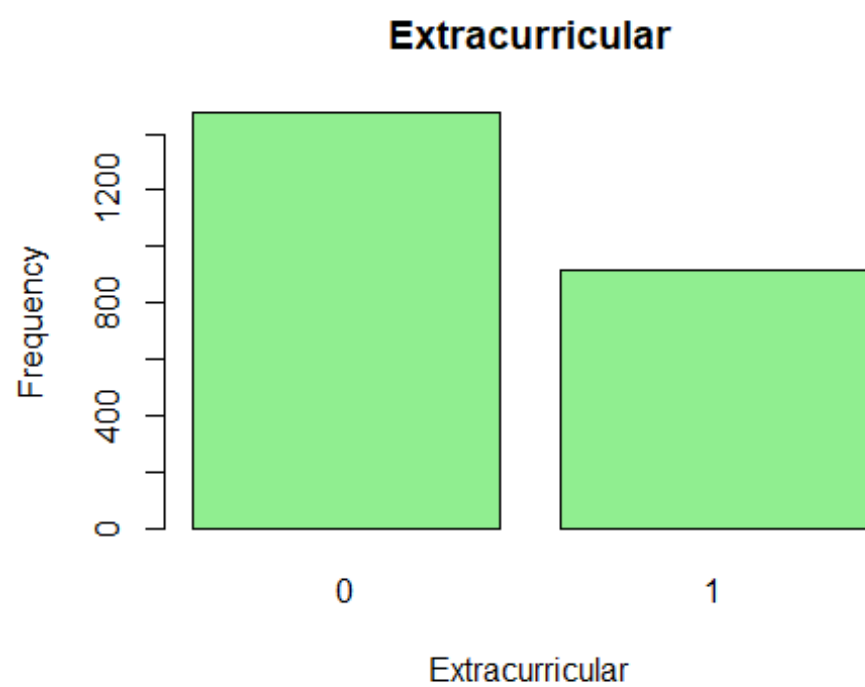


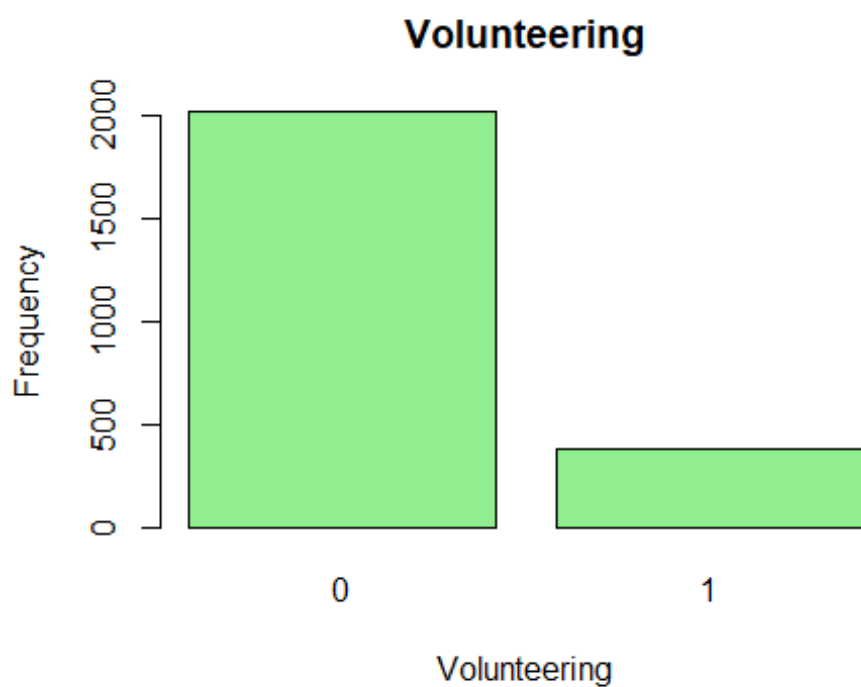
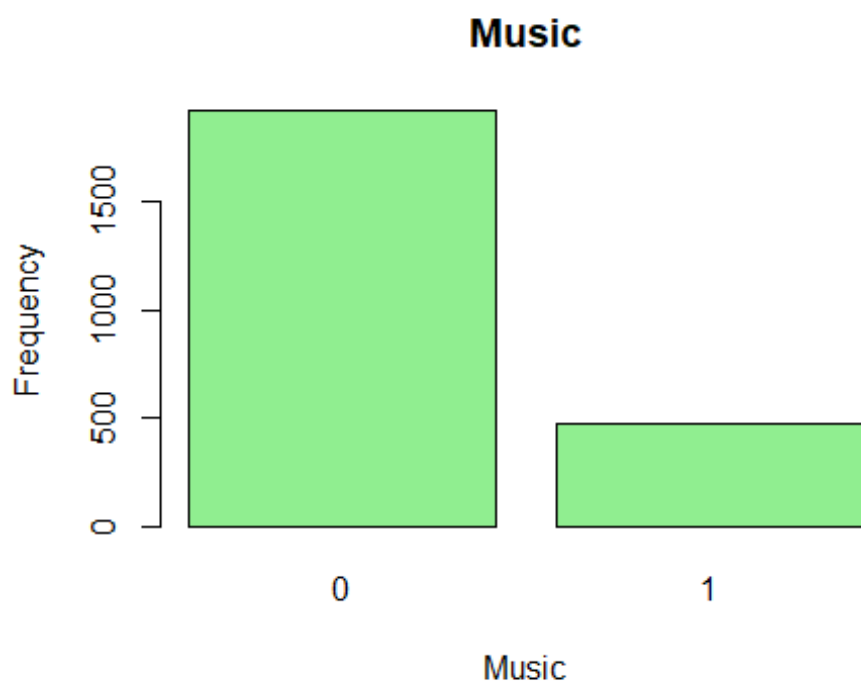
```
categorical=mydata[, -c(5,6,13)]
for(i in 1:ncol(categorical))
{
  a=table(categorical[,i])
  barplot(a, main=colnames(categorical)[i],
          xlab=colnames(categorical)[i],
          ylab="Frequency", col="lightgreen")
}
```








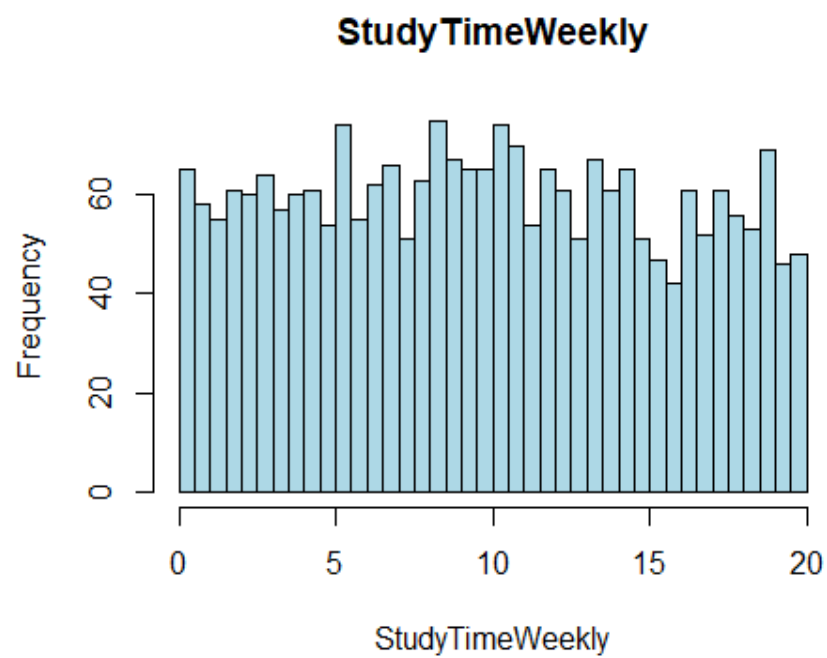





```
numerical=mydata0[,c(6,7,14)]
for(i in 1:ncol(numerical))
{
  hist(numerical[,i],main=colnames(numerical)[i],nclass=30,
```



```
xlab=colnames(numerical)[i],  
ylab="Frequency",col="lightblue")
```





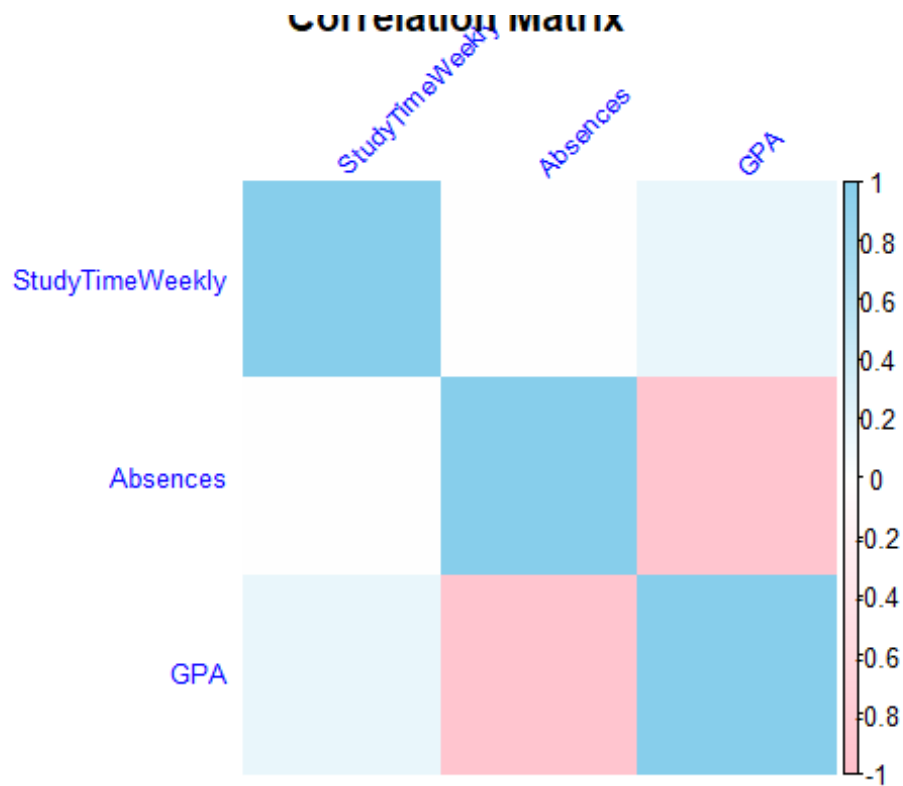
```
library(corrplot)

## Warning: package 'corrplot' was built under R version 4.4.2

## corrplot 0.95 loaded


correlation_matrix=cor(numerical, use = "complete.obs")
corrplot(correlation_matrix, method = "color",
         col = colorRampPalette(c("pink","white", "skyblue"))(200),
         tl.col = "blue",tl.cex = 0.8, tl.srt = 45, title = "Correlation Matrix")

## Warning in ind1:ind2: numerical expression has 2 elements: only the first used
```




```
x=mydata0[,-c(1,6,7,15)]
for(i in 1:(ncol(x)-1))
{
  boxplot(GPA~x[,i],data=x,main=paste("Boxplot of GPA vs",colnames(x)[i]),
    xlab=colnames(x)[i],ylab="GPA",
    las = 2,cex.names = 0.6)
}
```

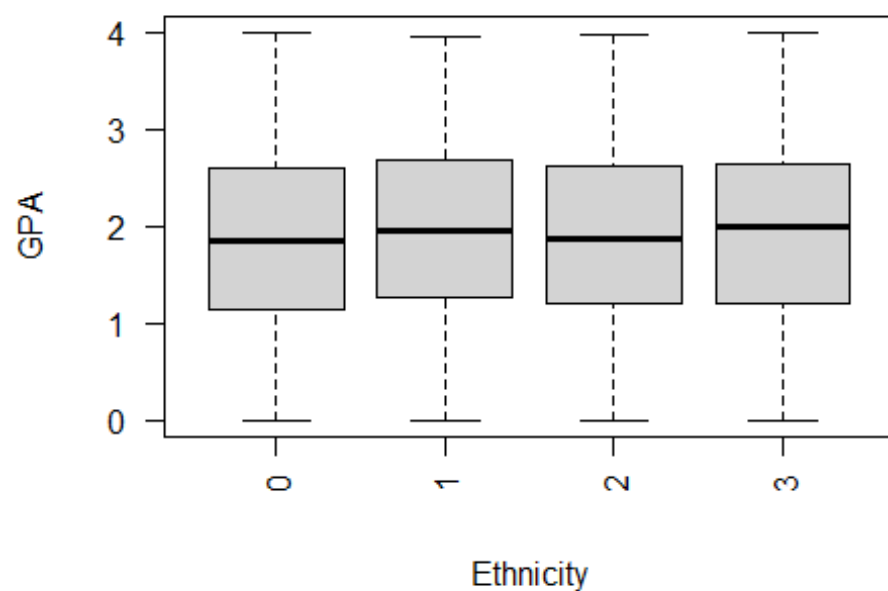
**Boxplot of GPA vs Age**




**Boxplot of GPA vs Gender**



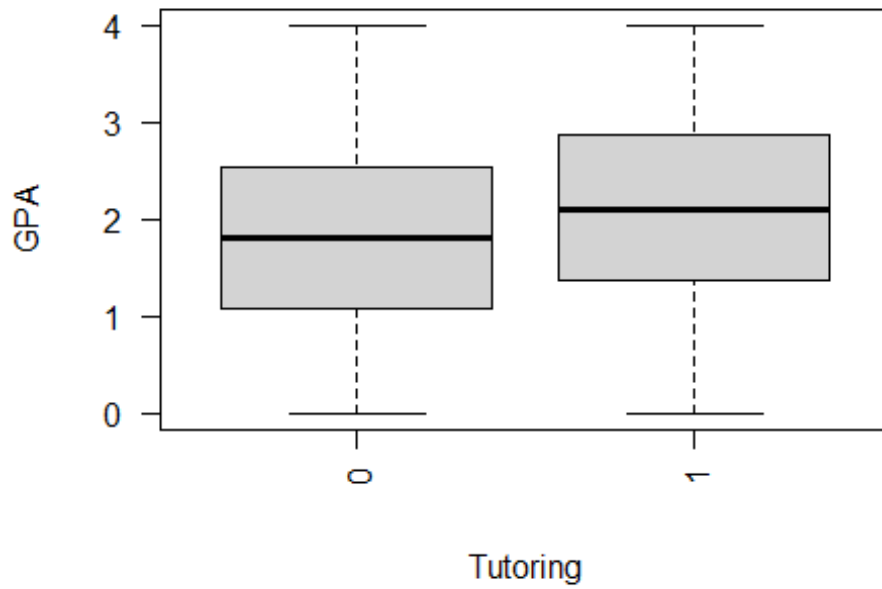
**Boxplot of GPA vs Ethnicity**




**Boxplot of GPA vs ParentalEducation**



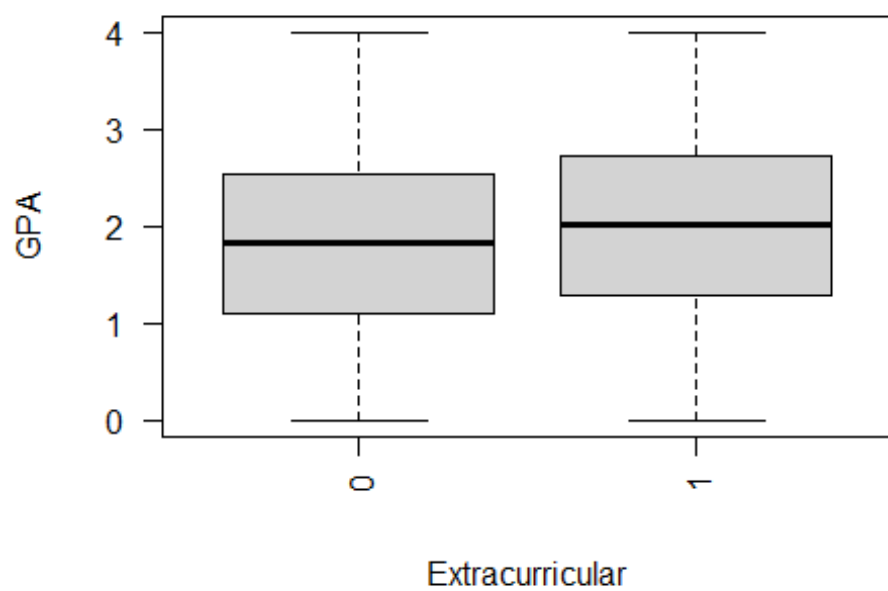
**Boxplot of GPA vs Tutoring**




**Boxplot of GPA vs ParentalSupport**



**Boxplot of GPA vs Extracurricular**




**Boxplot of GPA vs Sports**

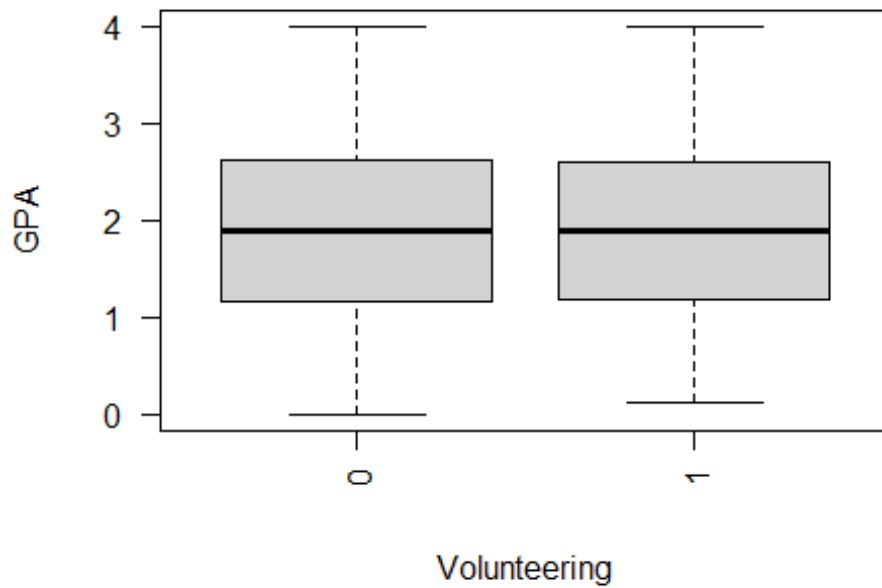




**Boxplot of GPA vs Music**



**Boxplot of GPA vs Volunteering**



# Model Fitting

## 1. Multinomial Logistic Regression

Data Preparation: Creating dummy variables and splitting dataset as train and test.

```
data1=mydata

data1$eth_caucasian=ifelse(data1$Ethnicity==0,1,0)
data1$eth_African_American=ifelse(data1$Ethnicity==1,1,0)
data1$eth_Asian=ifelse(data1$Ethnicity==2,1,0)

data1$par_ed_High_school=ifelse(data1$ParentalEducation==1,1,0)
data1$par_ed_college=ifelse(data1$ParentalEducation==2,1,0)
data1$par_ed_bachelor=ifelse(data1$ParentalEducation==3,1,0)
data1$par_ed_higher=ifelse(data1$ParentalEducation==4,1,0)

data1$Par_sup_low=ifelse(data1$ParentalSupport==1,1,0)
data1$Par_sup_moderate=ifelse(data1$ParentalSupport==2,1,0)
data1$Par_sup_high=ifelse(data1$ParentalSupport==3,1,0)
data1$Par_sup_very_high=ifelse(data1$ParentalSupport==4,1,0)

data2=data1[, -c(3,4,8)]

idx=sample(nrow(data2), size = 0.7*nrow(data2), replace = FALSE)
train=data2[idx,]
test=data2[-idx,]
```

Fitting model on train data set.

```
library(nnet)
logistic_model = multinom(GradeClass ~ ., data =data.frame(train))

## # weights: 110 (84 variable)
## initial value 2694.199065
## iter 10 value 1722.430367
## iter 20 value 1521.268520
## iter 30 value 1379.307227
## iter 40 value 1352.495807
## iter 50 value 1342.955388
## iter 60 value 1340.478248
## iter 70 value 1339.823053
## iter 80 value 1339.676179
## iter 90 value 1339.643717
## final value 1339.643635
## converged

summary(logistic_model)
```

```
## Call:
## multinom(formula = GradeClass ~ ., data = data.frame(train))
##
## Coefficients:
## (Intercept)          Age          Gender StudyTimeWeekly    Absences    Tutoring
## B      4.002384 -0.084482272 0.09350458      -0.03530022 -0.04951313 -0.44952
## 46
## C      4.791461 -0.010500704 0.16533453      -0.09164578  0.11515031 -1.03544
## 36
## D      4.546972 -0.051529922 0.13797502      -0.13920546  0.30682392 -1.56977
## 40
## F      2.051251 -0.009533597 0.13112417      -0.23082973  0.63605777 -2.19351
## 90
## Extracurricular    Sports          Music Volunteering eth_caucasian
## B      -0.07108413  0.1697829  0.2617197  -0.09002078  -0.7862314
## C      -0.25228041 -0.1874577 -0.2074666  -0.14942263  -0.8118651
## D      -0.79481638 -0.3136178 -0.4355552  -0.14661753  -0.8209159
## F      -1.12212270 -0.7580979 -0.5016813   0.02597506  -0.3164758
## eth_African_American eth_Asian par_ed_High_school par_ed_college
## B      -0.7693074 -1.4027435      0.032084705  0.08827789
## C      -0.7683193 -1.2126765      0.003273188  -0.03481707
## D      -0.5564450 -0.9027420      0.350282129  0.03850479
## F      -0.4231849 -0.6325834      -0.044366137 -0.20864999
## par_ed_bachelor par_ed_higher Par_sup_low Par_sup_moderate Par_sup_high
## B      -0.33019108  0.8699894  0.3384141  0.7484593  -0.3003643
## C      -0.26436657  0.6646912 -0.6398722  -0.7408678  -2.1904080
## D      -0.04156643  1.2895607 -0.4163825  -0.9736310  -2.5452884
## F      -0.26916788  1.3801192 -1.1272261  -2.1247821  -4.2760893
## Par_sup_very_high
## B      -0.9464278
## C      -2.8299110
## D      -3.1116559
## F      -5.2919763
##
## Std. Errors:
## (Intercept)          Age          Gender StudyTimeWeekly    Absences    Tutoring
## B      2.410577 0.1278449 0.2851014      0.02653595 0.03308449 0.2887667
## C      2.329792 0.1247437 0.2797657      0.02625735 0.03134548 0.2875915
## D      2.414245 0.1299794 0.2927455      0.02774533 0.03369122 0.3085216
## F      2.603644 0.1413947 0.3180574      0.03025997 0.03851930 0.3382356
## Extracurricular    Sports          Music Volunteering eth_caucasian
## B      0.2878931 0.3111430 0.3404571  0.3768517  0.5971815
## C      0.2825833 0.3086548 0.3408288  0.3727588  0.5914160
## D      0.2990582 0.3206466 0.3602217  0.3899449  0.6128225
## F      0.3257401 0.3478853 0.3915714  0.4252097  0.6601993
## eth_African_American eth_Asian par_ed_High_school par_ed_college
## B      0.6455340 0.6392624      0.5123817  0.4933839
## C      0.6387088 0.6251921      0.4956124  0.4786742
## D      0.6610988 0.6460518      0.5203162  0.5061917
```

```
## F          0.7142868 0.6998570          0.5591249          0.5417717
## par_ed_bachelor par_ed_higher Par_sup_low Par_sup_moderate Par_sup_high
## B          0.5738513          1.180251          0.9700283          0.9098058          0.8818479
## C          0.5530251          1.158186          0.8763875          0.8204472          0.7918183
## D          0.5818256          1.162634          0.9012985          0.8490933          0.8223991
## F          0.6258513          1.198893          0.9335442          0.8822694          0.8597359
## Par_sup_very_high
## B          0.9117061
## C          0.8266716
## D          0.8612576
## F          0.9127389
##
## Residual Deviance: 2679.287
## AIC: 2847.287
```

The objective of multinomial logistic regression model is to minimize negative log likelihood function which is attained at value 1410.87

```
training_pred=predict(logistic_model, newdata = train)
training_conf_matrix=table(Predicted = training_pred, Actual = train$GradeClass)
print(training_conf_matrix)
```

```
##          Actual
## Predicted  A   B   C   D   F
##          A   9   9   5   1   0
##          B  40 126  30   5   2
##          C   8  38 186  56   4
##          D   6   4  48 161  32
##          F  12  16  13  72 791
```

```
training_accuracy =sum(diag(training_conf_matrix)) / sum(training_conf_matrix)
```

```
training_error_rate=1-training_accuracy
```

```
print(paste("training Error Rate:", round(training_error_rate*100,2),"%"))
```

```
## [1] "training Error Rate: 23.95 %"
```

```
#training_accuracy = mean(training_pred == train$GradeClass)
```

```
cat("Training_Accuracy:", training_accuracy, "\n")
```

```
## Training_Accuracy: 0.760454
```

```
# Make predictions
```

```
testing_pred=predict(logistic_model, newdata = test)
```

```
testing_conf_matrix=table(Predicted = testing_pred, Actual = test$GradeClass)
```

```
print(testing_conf_matrix)
```

```
##          Actual
## Predicted  A   B   C   D   F
##          A   0   5   2   1   0
##          B  21  34  12   1   3
```

```
##           C    3   18   66   22    2
##           D    3    5   20   57   20
##           F    5   14    9   38  357

testing_accuracy = sum(diag(testing_conf_matrix)) / sum(testing_conf_matrix)
testing_error_rate = 1 - testing_accuracy
print(paste("Testing Error Rate:", round(testing_error_rate*100,2), "%"))

## [1] "Testing Error Rate: 28.41 %"

#testing_accuracy = mean(testing_pred == test$GradeClass)
cat("Testing_Accuracy:", testing_accuracy, "\n")

## Testing_Accuracy: 0.7158774
```

Variable selection:

```
library(MASS)
step=stepAIC(logistic_model,direction="both")

summary(step)$call

## multinom(formula = GradeClass ~ StudyTimeWeekly + Absences +
##           Tutoring + Extracurricular + Sports + Music + Par_sup_moderate +
##           Par_sup_high + Par_sup_very_high, data = data.frame(train))

new_logistic=multinom(formula = GradeClass ~ StudyTimeWeekly + Absences +
  Tutoring + Extracurricular + Sports + Music + par_ed_higher +
  Par_sup_low + Par_sup_moderate + Par_sup_high + Par_sup_very_high,
  data = data.frame(train))

## # weights:  65 (48 variable)
## initial  value 2694.199065
## iter   10 value 1591.913859
## iter   20 value 1373.516149
## iter   30 value 1352.901289
## iter   40 value 1349.899197
## iter   50 value 1349.208017
## final   value 1349.194235
## converged

summary(new_logistic)

## Call:
## multinom(formula = GradeClass ~ StudyTimeWeekly + Absences +
##           Tutoring + Extracurricular + Sports + Music + par_ed_higher +
##           Par_sup_low + Par_sup_moderate + Par_sup_high + Par_sup_very_high,
##           data = data.frame(train))
##
## Coefficients:
## (Intercept) StudyTimeWeekly Absences Tutoring Extracurricular Sp
orts
```

```

## B      1.580843      -0.02780384 -0.04833114 -0.449087      -0.07368919  0.108
0934
## C      3.603522      -0.08462679  0.11611936 -1.032889      -0.24925407 -0.260
5872
## D      2.984008      -0.13404452  0.30641826 -1.557989      -0.77318497 -0.373
5747
## F      1.259938      -0.22623241  0.63495833 -2.184170      -1.11666009 -0.814
5800
##          Music par_ed_higher Par_sup_low Par_sup_moderate Par_sup_high
## B  0.2742091      0.8763296  0.4542169      0.8384799  -0.1640848
## C -0.2119239      0.6776981 -0.5225228      -0.6590190  -2.0686972
## D -0.4344324      1.1252143 -0.3154435      -0.8974957  -2.4355269
## F -0.5119093      1.4275320 -0.9682236      -2.0106828  -4.1347214
## Par_sup_very_high
## B      -0.7770625
## C      -2.6619055
## D      -2.9579594
## F      -5.0779217
##
## Std. Errors:
## (Intercept) StudyTimeWeekly Absences Tutoring Extracurricular Sports
## B  0.9318091      0.02601207 0.03276366 0.2842735      0.2845906 0.30642
62
## C  0.8413955      0.02579585 0.03120948 0.2842015      0.2798943 0.30438
22
## D  0.8658642      0.02733678 0.03358725 0.3051598      0.2965241 0.31655
37
## F  0.9075125      0.02987222 0.03838097 0.3349066      0.3229807 0.34380
04
##          Music par_ed_higher Par_sup_low Par_sup_moderate Par_sup_high
## B 0.3354433      1.092915  0.9594989      0.9002224  0.8699449
## C 0.3373349      1.076004  0.8643480      0.8090939  0.7783818
## D 0.3572102      1.070625  0.8888717      0.8370161  0.8088038
## F 0.3888730      1.098757  0.9206762      0.8697888  0.8460421
## Par_sup_very_high
## B      0.9005099
## C      0.8137335
## D      0.8481894
## F      0.8980391
##
## Residual Deviance: 2698.388
## AIC: 2794.388

training_pred=predict(new_logistic, newdata = train)
training_conf_matrix=table(Predicted = training_pred, Actual = train$GradeClass)
print(training_conf_matrix)

```

```

##           Actual
## Predicted  A    B    C    D    F
##           A    6    6    6    0    0
##           B   43  127   30    5    2
##           C    8   40  189   54    6
##           D    6    6   46  161   32
##           F   12   14   11   75  789

training_accuracy = sum(diag(training_conf_matrix)) / sum(training_conf_matrix)
training_error_rate = 1 - training_accuracy
print(paste("training Error Rate:", round(training_error_rate*100,2), "%"))

## [1] "training Error Rate: 24.01 %"

#training_accuracy = mean(training_pred == train$GradeClass)
cat("Training_Accuracy:", training_accuracy, "\n")

## Training_Accuracy: 0.7598566

# Make predictions
testing_pred = predict(new_logistic, newdata = test)
testing_conf_matrix = table(Predicted = testing_pred, Actual = test$GradeClass)
print(testing_conf_matrix)

##           Actual
## Predicted  A    B    C    D    F
##           A    0    2    2    1    0
##           B   20   37   12    1    3
##           C    4   18   66   20    2
##           D    3    5   20   62   16
##           F    5   14    9   35  361

testing_accuracy = sum(diag(testing_conf_matrix)) / sum(testing_conf_matrix)
testing_error_rate = 1 - testing_accuracy
print(paste("Testing Error Rate:", round(testing_error_rate*100,2), "%"))

## [1] "Testing Error Rate: 26.74 %"

#testing_accuracy = mean(testing_pred == test$GradeClass)
cat("Testing_Accuracy:", testing_accuracy, "\n")

## Testing_Accuracy: 0.7325905

summary(logistic_model)$AIC

## [1] 2847.287

summary(new_logistic)$AIC

## [1] 2794.388

```

Model	Training accuracy	Testing accuracy	AIC	Residual deviance
Logistic model	0.7605	0.7158	2847.287	2679.287
New Logistic model	0.7598	0.7325	2794.388	2698.388

Conclusion:

- AIC value of new model is less than model fitted using all variables. Hence the new model is better.
- Training and testing accuracy also increases for new model.



## 2. K- nearest neighbor classifier

data preparation and splitting

```
mydata=mydata0[,-c(1,14)]
idx=sample(nrow(mydata), size = 0.7*nrow(mydata), replace = FALSE)
train=mydata[idx,]
test=mydata[-idx,]
dim(train)


## [1] 1674    13

y_train=train[,13]
y_test=test[,13]
x_train=train[,-13]
x_test=test[,-13]
```

fitting KNN model for different values of k

```
library(class)
k=1:100
knn_accuracy=knn_error_rate=c()
for(kn in 1:100){
  knn_pred=knn(train = x_train,test = x_test,
               cl=train$GradeClass,k=kn)
  conf_matrix=table(Predicted = knn_pred, Actual = test$GradeClass)
  knn_accuracy[kn]=sum(diag(conf_matrix)) / sum(conf_matrix)
  knn_error_rate[kn]=1-knn_accuracy[kn]
}
plot(k,knn_error_rate,type='l',main="Plot of error rate vs k")
```

**Plot of error rate vs k**



Tuning for best value of k

```
library(e1071)
kn=1:100
tune_knn=tune.knn(train[, -13], as.factor(train[, 13]), k=kn)
summary(tune_knn)

##
## Parameter tuning of 'knn.wrapper':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   k
## 22
##
## - best performance: 0.2943201
##
## - Detailed performance results:
##      k      error dispersion
## 1    1 0.3879480 0.02753620
## 2    2 0.3871130 0.02451922
## 3    3 0.3411280 0.03593381
## 4    4 0.3285983 0.04155377
## 5    5 0.3135617 0.02625928
## 6    6 0.3177563 0.03969032
## 7    7 0.3156450 0.03362630
```

## 8	8	0.3160739	0.03181845
## 9	9	0.3160652	0.03037944
## 10	10	0.3093759	0.02628831
## 11	11	0.3093672	0.02400584
## 12	12	0.3139592	0.02362528
## 13	13	0.3152092	0.02633696
## 14	14	0.3101953	0.02391954
## 15	15	0.3118724	0.01659872
## 16	16	0.3068654	0.02626206
## 17	17	0.3051883	0.03155836
## 18	18	0.3072891	0.03257256
## 19	19	0.3022681	0.02682674
## 20	20	0.3081224	0.02806074
## 21	21	0.3014296	0.02713119
## 22	22	<b>0.2943201</b>	<b>0.02745624</b>
## 23	23	0.2943201	0.02766036
## 24	24	0.3022577	0.02859997
## 25	25	0.3005927	0.03241555
## 26	26	0.3018445	0.03297950
## 27	27	0.2993358	0.02567518
## 28	28	0.2993515	0.03299333
## 29	29	0.2989331	0.03029602
## 30	30	0.3006084	0.02814749
## 31	31	0.3014400	0.02843462
## 32	32	0.3014400	0.02870696
## 33	33	0.3043689	0.03468679
## 34	34	0.3043602	0.03148821
## 35	35	0.3077040	0.02979397
## 36	36	0.3097960	0.03637566
## 37	37	0.3064435	0.03832237
## 38	38	0.3093759	0.03610164
## 39	39	0.3102075	0.03566634
## 40	40	0.3072821	0.03245561
## 41	41	0.3093741	0.03055379
## 42	42	0.3072856	0.03139352
## 43	43	0.3047751	0.03272813
## 44	44	0.3047751	0.03052203
## 45	45	0.3064470	0.03386608
## 46	46	0.3051953	0.03289509
## 47	47	0.3102127	0.03270456
## 48	48	0.3077092	0.03355284
## 49	49	0.3056172	0.03378178
## 50	50	0.3068706	0.03683568
## 51	51	0.3093863	0.03574668
## 52	52	0.3052057	0.03537916
## 53	53	0.3110600	0.03647930
## 54	54	0.3102249	0.03642444
## 55	55	0.3089714	0.03806960
## 56	56	0.3089679	0.04079832
## 57	57	0.3072960	0.03695758

```
## 58 58 0.3056224 0.03782423
## 59 59 0.3068741 0.03573872
## 60 60 0.3072943 0.03698451
## 61 61 0.3068776 0.03741497
## 62 62 0.3068776 0.03657368
## 63 63 0.3077127 0.04006486
## 64 64 0.3018602 0.03551804
## 65 65 0.3031189 0.03569613
## 66 66 0.3081311 0.03713742
## 67 67 0.3110617 0.03425082
## 68 68 0.3064575 0.03503136
## 69 69 0.3072960 0.03673852
## 70 70 0.3085495 0.03791314
## 71 71 0.3064627 0.03695978
## 72 72 0.3035356 0.03764228
## 73 73 0.3089662 0.03421975
## 74 74 0.3093846 0.03431098
## 75 75 0.3089662 0.03329782
## 76 76 0.3085478 0.03383634
## 77 77 0.3064557 0.03290158
## 78 78 0.3043654 0.03186833
## 79 79 0.3085478 0.03337327
## 80 80 0.3060321 0.03312588
## 81 81 0.3047786 0.03287091
## 82 82 0.3068724 0.03271590
## 83 83 0.3081259 0.03415159
## 84 84 0.3060356 0.03367267
## 85 85 0.3077075 0.03443101
## 86 86 0.3077040 0.03433289
## 87 87 0.3072856 0.03394396
## 88 88 0.3081224 0.03587928
## 89 89 0.3081241 0.03717707
## 90 90 0.3097978 0.03839372
## 91 91 0.3102144 0.03834502
## 92 92 0.3089575 0.03830882
## 93 93 0.3077022 0.03680184
## 94 94 0.3093759 0.03405876
## 95 95 0.3077040 0.03379030
## 96 96 0.3077022 0.03405667
## 97 97 0.3072856 0.03563075
## 98 98 0.3085391 0.03444624
## 99 99 0.3110513 0.03210595
## 100 100 0.3093724 0.03372986

paste("Best value of k is", tune_knn$best.parameters)

## [1] "Best value of k is 22"
```

Prediction for test data

```

knn_pred_test=knn(train = train[, -13], test = test[, -13],
                  cl=train$GradeClass, k=tune_knn$best.parameters$k)
conf_matrix_test=table(Predicted = knn_pred_test, Actual = test$GradeClass)
conf_matrix_test

##           Actual
## Predicted   A    B    C    D    F
##           A     3     2     0     0     0
##           B    13    39    20     2     1
##           C     7    38    66    29     4
##           D     1     2    23    63    16
##           F     5    13     9    40   322

knn_test_accuracy=sum(diag(conf_matrix_test)) / sum(conf_matrix_test)
knn_test_error_rate=1-knn_test_accuracy
knn_test_error_rate

## [1] 0.3133705

```

Model	Best value of K	Testing Accuracy
KNN	22	0.6866

### 3. Support Vector Machine

Finding optimum value of parameters (cost, gamma) for radial kernel

```
library(e1071)
### Tuning for radial kernel
set.seed(1)
tune.out_radial=tune(svm, as.factor(GradeClass)~., data = train,
                     kernel = "radial",
                     ranges = list(
                       cost = c(0.1,0.5, 1,5, 10),
                       gamma = c(0.5,0.75,1,2)
                     )
                   )
summary(tune.out_radial)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##     5    0.5
##
## - best performance: 0.4209571
##
## - Detailed performance results:
##   cost gamma    error dispersion
## 1  0.1  0.50 0.4937134 0.02800566
## 2  0.5  0.50 0.4698815 0.03363581
## 3  1.0  0.50 0.4238790 0.03582883
## 4  5.0  0.50 0.4209571 0.03187642
## 5 10.0  0.50 0.4217974 0.02836011
## 6  0.1  0.75 0.4937134 0.02800566
## 7  0.5  0.75 0.4928783 0.02739753
## 8  1.0  0.75 0.4564906 0.03565685
## 9  5.0  0.75 0.4410303 0.02753332
## 10 10.0 0.75 0.4414487 0.02718091
## 11 0.1  1.00 0.4937134 0.02800566
## 12 0.5  1.00 0.4937134 0.02800566
## 13 1.0  1.00 0.4748937 0.03089951
## 14 5.0  1.00 0.4602545 0.03191472
## 15 10.0 1.00 0.4602545 0.03129929
## 16 0.1  2.00 0.4937134 0.02800566
## 17 0.5  2.00 0.4937134 0.02800566
## 18 1.0  2.00 0.4899442 0.02885234
```

```

## 19  5.0  2.00 0.4857636 0.02831151
## 20 10.0  2.00 0.4857636 0.02831151

tune.out_radial$best.parameters

##    cost gamma
## 4     5    0.5

tune.out_radial$best.performance

## [1] 0.4209571

### Tuning for linear kernel
tune.out_linear=tune(svm, as.factor(GradeClass)~., data = train,
                      kernel = "linear",
                      ranges = list(cost = c(0.1,0.5, 1,5, 10))
                      )

summary(tune.out_linear)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.2399808
##
## - Detailed performance results:
##   cost      error dispersion
## 1  0.1 0.2462517 0.03300885
## 2  0.5 0.2429079 0.03763766
## 3  1.0 0.2399808 0.03322843
## 4  5.0 0.2433264 0.03047065
## 5 10.0 0.2416527 0.03020263

tune.out_linear$best.parameters

##    cost
## 3     1

tune.out_linear$best.performance

## [1] 0.2399808

### plots
plot(subset(tune.out_radial$performances,tune.out_radial$performances$gamma==
0.5)$cost,subset(tune.out_radial$performances,tune.out_radial$performances$gamma==0.5)$error,type="l",ylim=c(0.2,0.8),col="red",lwd=1.5,main="cost vs error")

```

```

r for SVM",xlab="cost",ylab="error")

lines(subset(tune.out_radial$performances,tune.out_radial$performances$gamma=
=0.75)$cost,subset(tune.out_radial$performances,tune.out_radial$performances$
gamma==0.75)$error,type="l",col="blue",lwd=1.5)

lines(subset(tune.out_radial$performances,tune.out_radial$performances$gamma=
=1)$cost,subset(tune.out_radial$performances,tune.out_radial$performances$gam
ma==1)$error,type="l",col="purple",lwd=1.5)

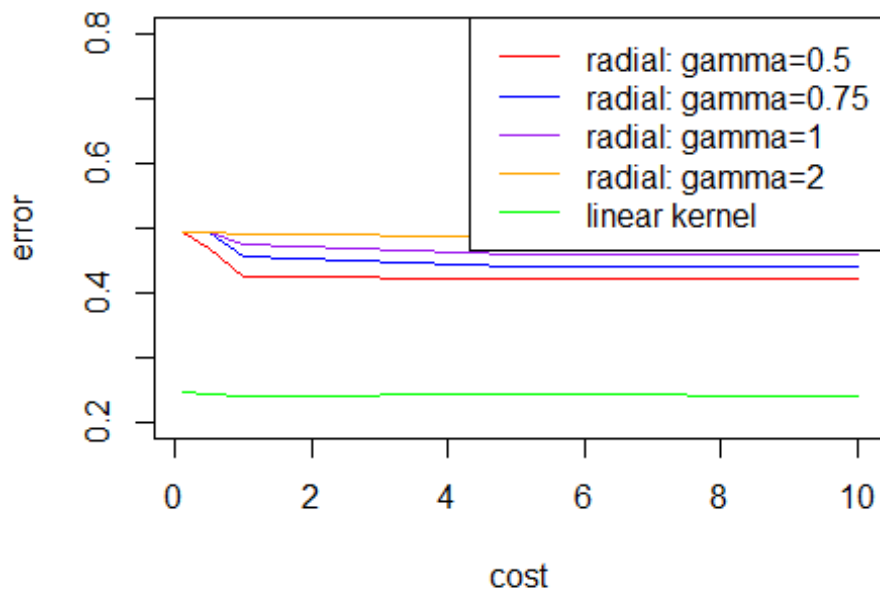
lines(subset(tune.out_radial$performances,tune.out_radial$performances$gamma=
=2)$cost,subset(tune.out_radial$performances,tune.out_radial$performances$gam
ma==2)$error,type="l",col="orange",lwd=1.5)

lines(tune.out_linear$performances$cost,tune.out_linear$performances$error,ty
pe="l",col="green",lwd=1.5)

legend("topright",c("radial: gamma=0.5","radial: gamma=0.75","radial: gamma=1
","radial: gamma=2","linear kernel"),col=c("red","blue","purple","orange","gr
een"),lty=rep(1,4))

```

### cost vs error for SVM





Kernel	Best Cost	Best Gamma	Performance (Error)
Radial	5	0.5	0.4209
Linear	1	-	0.23998

```

### svm for validation
svmfit=svm(as.factor(GradeClass)~., data = train,
           kernel = "linear",
           cost = 1)
summary(svmfit)

##
## Call:
## svm(formula = as.factor(GradeClass) ~ ., data = train, kernel = "linear",
##      cost = 1)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:   1
##
## Number of Support Vectors:  962
##
## ( 169 78 280 266 169 )
##
##
## Number of Classes:  5
##
## Levels:
##  A B C D F

train_err=1-sum(diag(table(svmfit$fitted, train$GradeClass)))/sum(table(svmfi
t$fitted,train$GradeClass))
test_conf_matrix=table(predict(svmfit,test),test$GradeClass)
test_err=1-sum(diag(test_conf_matrix))/sum(test_conf_matrix)
paste("training error:",train_err)

## [1] "training error: 0.228793309438471"

paste("testing error:",test_err)

## [1] "testing error: 0.268802228412256"

```

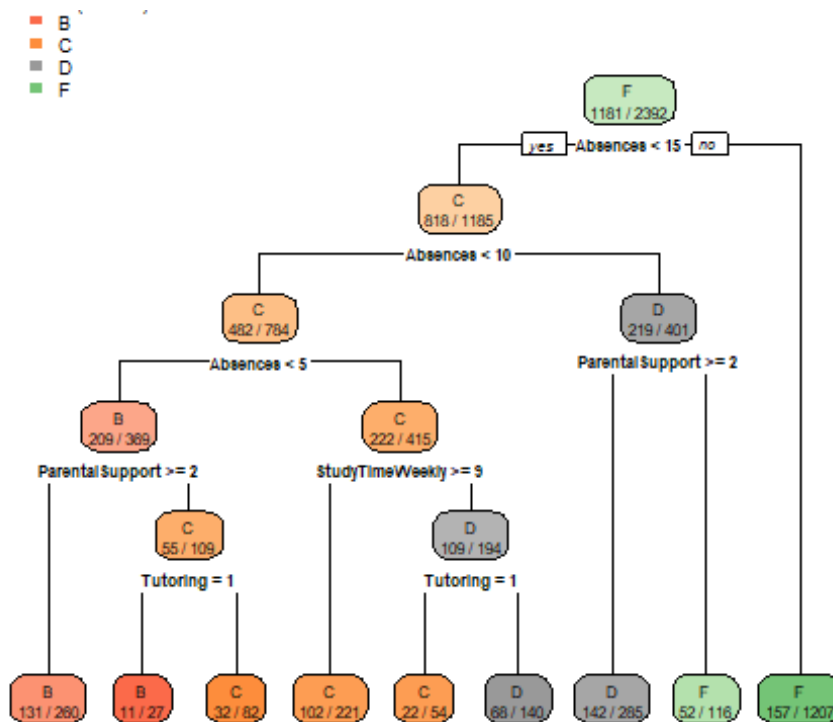
## 4. Decision Tree Classifier

Fitting a decision tree

```
library(tree)

library(rpart)
library(rpart.plot)

tree=rpart(as.factor(mydata$GradeClass)~.,data=mydata)
rpart.plot(tree,extra=3)
```



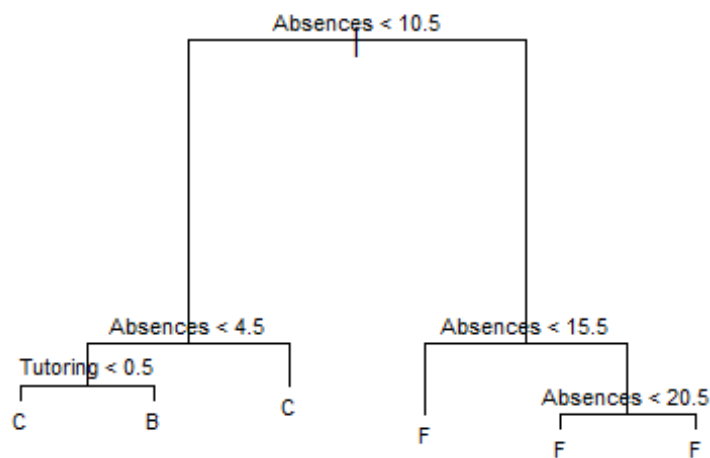
```
print(tree)

## n= 2392
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 2392 1181 F (0.045 0.11 0.16 0.17 0.51)
##    2) Absences< 14.5 1185 818 C (0.077 0.21 0.31 0.27 0.14)
##      4) Absences< 9.5 784 482 C (0.11 0.29 0.39 0.18 0.033)
##        8) Absences< 4.5 369 209 B (0.19 0.43 0.3 0.06 0.019)
##          16) ParentalSupport>=1.5 260 131 B (0.23 0.5 0.21 0.038 0.019) *
##            17) ParentalSupport< 1.5 109 55 C (0.092 0.28 0.5 0.11 0.018)
##              34) Tutoring>=0.5 27 11 B (0.26 0.59 0.15 0 0) *
##              35) Tutoring< 0.5 82 32 C (0.037 0.18 0.61 0.15 0.024) *
```

```
##      9) Absences>=4.5 415  222 C (0.031 0.17 0.47 0.29 0.046)
##      18) StudyTimeWeekly>=8.957582 221  102 C (0.036 0.24 0.54 0.15 0.0
27) *
##      19) StudyTimeWeekly< 8.957582 194  109 D (0.026 0.088 0.38 0.44 0.
067)
##      38) Tutoring>=0.5 54  22 C (0.019 0.15 0.59 0.24 0) *
##      39) Tutoring< 0.5 140  68 D (0.029 0.064 0.3 0.51 0.093) *
##      5) Absences>=9.5 401  219 D (0.017 0.03 0.16 0.45 0.34)
##      10) ParentalSupport>=1.5 285  142 D (0.018 0.039 0.19 0.5 0.25) *
##      11) ParentalSupport< 1.5 116  52 F (0.017 0.0086 0.086 0.34 0.55) *
##      3) Absences>=14.5 1207  157 F (0.013 0.022 0.02 0.075 0.87) *
```

fitting decision tree on train data

```
#Training tree model
train_tree=tree(as.factor(train$GradeClass)~.,data = train)
plot(train_tree)
text(train_tree,cex=0.7)
```



prediction and accuracy:

```
#Tree prediction on test
pr_test_tree=predict(train_tree,newdata=test,type = 'class')
pr_train_tree=predict(train_tree,newdata=train,type="class")

#confusion matrix
a=table(pr_test_tree,test$GradeClass)
```

```
test_accuracy=sum(diag(a))/sum(a)
paste("test accuracy:",test_accuracy)

## [1] "test accuracy: 0.61142061281337"

b=table(pr_train_tree,train$GradeClass)
train_accuracy=sum(diag(b))/sum(b)
paste("train accuracy:",train_accuracy)

## [1] "train accuracy: 0.663679808841099"
```

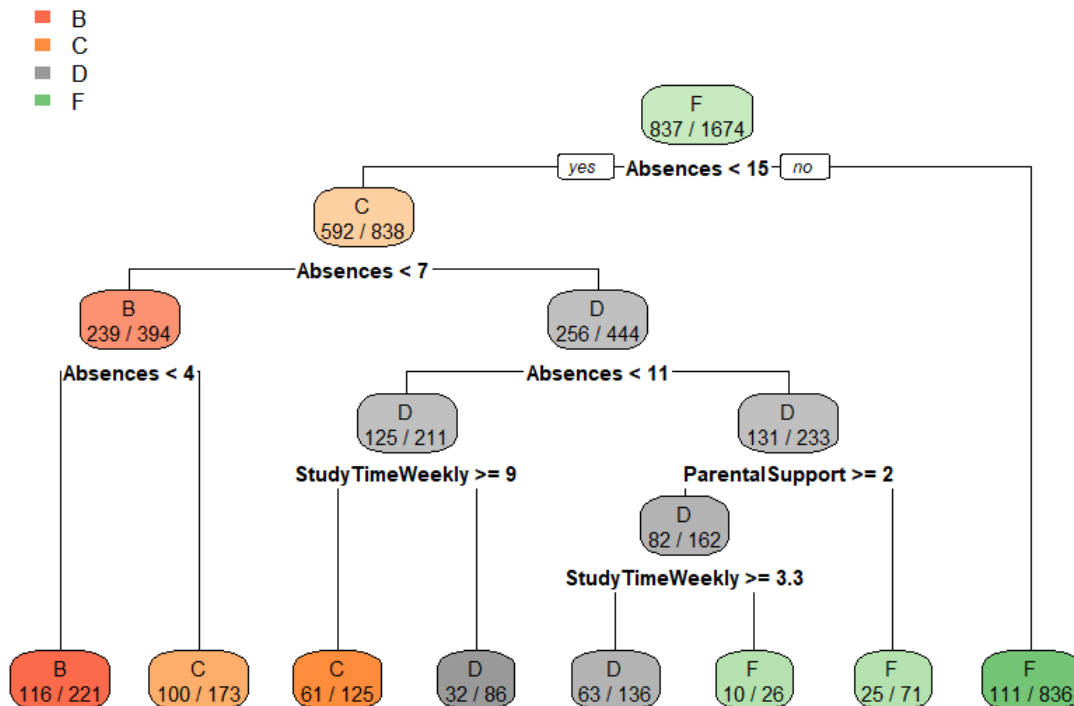
Fitting models using gini index and entropy as splitting criteria.

```
tree_model_gini=rpart(as.factor(mydata$GradeClass) ~ ., data = train, parms =
list(split = "gini"))

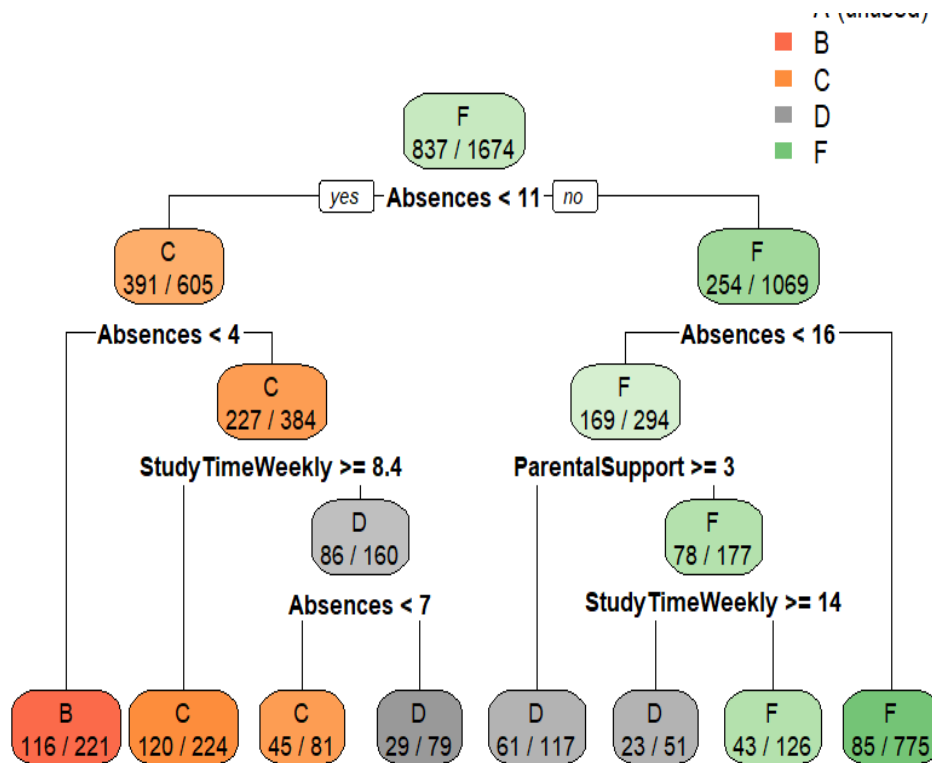
tree_model_entropy=rpart(as.factor(GradeClass) ~ ., data = train, parms = lis
t(split = "information"))
```

Pruning tree

```
prune.train_tree_gini=prune(tree_model_gini,cp=0.010000)
rpart.plot(prune.train_tree_gini,extra=3)
```



```
prune.train_tree_entropy=prune(tree_model_entropy,cp=0.010000)
rpart.plot(prune.train_tree_entropy,extra=3)
```



```
prune.test_tree_gini=predict(prune.train_tree_gini,test,type = 'class')
confusion_matrix_gini=as.matrix(table(prune.test_tree_gini,test$GradeClass))
accuracy_gini=sum(diag(confusion_matrix_gini))/sum(confusion_matrix_gini);acc
uracy_gini
```

```
## [1] 0.6824513
```

```
prune.test_tree_entropy=predict(prune.train_tree_entropy,test,type = 'class')
confusion_matrix_entropy=as.matrix(table(prune.test_tree_entropy,test$GradeCl
ass))
accuracy_entropy=sum(diag(confusion_matrix_entropy))/sum(confusion_matrix_ent
ropy);accuracy_entropy
```

```
## [1] 0.6713092
```

Splitting criterion	Cp for pruning	Accuracy
Gini index	0.01	0.6824
Entropy	0.01	0.6712

## 5.Bagging

K fold cv for bagging.

```
library(caret)

library(randomForest)
k = 10
set.seed(123)
folds=createFolds(mydata$GradeClass, k = k, list = TRUE, returnTrain = FALSE)
train_error_rates=test_error_rates=c()
Mean_train_Error_Rate=Mean_test_Error_Rate=c()
oob=c()
nseq=seq(10,200,by=10)
# Perform k-fold cross-validation
for (p in 1:length(nseq))
{
  for(i in 1:k )
  {
    test_indices = folds[[i]]
    x_train_cv = mydata[-test_indices, ]
    x_test_cv = mydata[test_indices, ]
    bag = randomForest(as.factor(GradeClass) ~ ., data = x_train_cv,
                      mtry = ncol(x_train_cv) - 1, ntree = nseq[p])

    #train
    train_pred_bag=predict(bag, newdata = x_train_cv)
    train_conf_matrix_bag=table(x_train_cv$GradeClass, train_pred_bag)
    train_error_rate=1 -sum(diag(train_conf_matrix_bag)) / sum(train_conf_mat
rix_bag)
    train_error_rates=c(train_error_rates, train_error_rate)
    #test
    test_pred_bag=predict(bag, newdata = x_test_cv)
    test_conf_matrix_bag=table(x_test_cv$GradeClass, test_pred_bag)
    test_error_rate=1 -sum(diag(test_conf_matrix_bag)) / sum(test_conf_matrix
_bag)
    test_error_rates=c(test_error_rates, test_error_rate)

  }
  Mean_train_Error_Rate[p]= mean(train_error_rates)
  Mean_test_Error_Rate[p]= mean(test_error_rates)


  bag_oob=randomForest(as.factor(GradeClass) ~ ., data = mydata,
                      mtry = ncol(mydata) - 1, ntree = nseq[p])
  oob[p]=bag_oob$err.rate[bag_oob$ntree,1]
}

plot(nseq,Mean_train_Error_Rate,type="l",col="red",ylim=c(0,1),lwd=1.5,xlab="
```

```

Number of trees",ylab="Error")
lines(nseq,Mean_test_Error_Rate,type="l",col="black",lwd=1.5)
lines(nseq,oob,type="l",col="green",lwd=1.5)

```



To get optimal number of trees.

```

tune_bag=tune.randomForest(as.factor(GradeClass)~.,data=train,
                           ntree=seq(10,200,by=10),mtry=ncol(train)-1)
summary(tune_bag)

```

```

##
## Parameter tuning of 'randomForest':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   mtry ntree
##    12   190
##
## - best performance: 0.2742591
##
## - Detailed performance results:
##   mtry ntree   error dispersion
## 1    12    10 0.2993602 0.03424495
## 2    12    20 0.2847106 0.02783673
## 3    12    30 0.2813720 0.03362080
## 4    12    40 0.2847089 0.02897221

```

```
## 5      12      50 0.2822106 0.03433976
## 6      12      60 0.2817957 0.03486036
## 7      12      70 0.2792852 0.02869607
## 8      12      80 0.2863842 0.03180385
## 9      12      90 0.2788563 0.02329375
## 10     12     100 0.2830457 0.03464706
## 11     12     110 0.2817957 0.03445816
## 12     12     120 0.2838755 0.02940360
## 13     12     130 0.2801203 0.03322604
## 14     12     140 0.2830370 0.02820180
## 15     12     150 0.2817904 0.02787687
## 16     12     160 0.2780230 0.02518271
## 17     12     170 0.2809554 0.02784100
## 18     12     180 0.2859693 0.02119490
## 19     12     190 0.2742591 0.02855679
## 20     12     200 0.2801116 0.02468123
```

For validation

```
bag_final_train=randomForest(as.factor(GradeClass)~.,data=train,
                             mtry=12,ntree=tune_bag$best.parameters$ntree)
print(bag_final_train)

##
## Call:
## randomForest(formula = as.factor(GradeClass) ~ ., data = train,      mtry
## = 12, ntree = tune_bag$best.parameters$ntree)
##           Type of random forest: classification
##           Number of trees: 190
## No. of variables tried at each split: 12
##
##           OOB estimate of  error rate: 27.78%
## Confusion matrix:
##   A  B   C   D   F class.error
## A 22 32   4   7  13  0.71794872
## B 13 85  52   9  16  0.51428571
## C  0 28 164  66  15  0.39926740
## D  1  8  63 137  71  0.51071429
## F  2  3   6  56 801  0.07718894


test_pred_bag_final=predict(bag_final_train, newdata =test)
test_conf_matrix_bag_final=table(test$GradeClass, test_pred_bag_final)
test_error_rate_bag_final=1 -sum(diag(test_conf_matrix_bag_final)) / sum(test
_conf_matrix_bag_final)
test_error_rate_bag_final

## [1] 0.2980501

barplot(sort(importance(bag_final)[,1]),horiz=T,
        las=1,cex.names = 0.4,main = "Variable Importance Plot")
```



## Variable Importance Plot



Method	Best n tree	OOB error	Testing error
Bagging	190	0.2778	0.2980

## 6.Random Forest

For different number of trees, plot of error rate of random forest model with different mtry values.

```
k = 5
set.seed(123)
folds=createFolds(mydata$GradeClass, k = k, list = TRUE, returnTrain = FALSE)
train_error_rates_rf1=train_error_rates_rf2=train_error_rates_rf3=train_error
_rates_rf4=c()
test_error_rates_rf1=test_error_rates_rf2=test_error_rates_rf3=test_error_rat
es_rf4=c()
Mean_train_Error_Rate_rf1=Mean_train_Error_Rate_rf2=Mean_train_Error_Rate_rf3
=Mean_train_Error_Rate_rf4=c()
Mean_test_Error_Rate_rf1=Mean_test_Error_Rate_rf2=Mean_test_Error_Rate_rf3=Me
an_test_Error_Rate_rf4=c()
oob_rf1=oob_rf2=oob_rf3=oob_rf4=c()
nseq=seq(10,200,by=10)
# Perform k-fold cross-validation
for (p in 1:length(nseq))
{
  for(i in 1:k )
  {
    test_indices = folds[[i]]
    x_train_cv = mydata[-test_indices, ]
    x_test_cv = mydata[test_indices, ]

    **** model1
    rf1 = randomForest(as.factor(GradeClass) ~ ., data = x_train_cv,
                      mtry = (ncol(x_train_cv)-1)/2, ntree = nseq[p])

    #train
    train_pred_rf1=predict(rf1, newdata = x_train_cv)
    train_conf_matrix_rf1=table(x_train_cv$GradeClass, train_pred_rf1)
    train_error_rate_rf1=1 -sum(diag(train_conf_matrix_rf1)) / sum(train_conf
_matrix_rf1)
    train_error_rates_rf1=c(train_error_rates_rf1, train_error_rate_rf1)
    #test
    test_pred_rf1=predict(rf1, newdata = x_test_cv)
    test_conf_matrix_rf1=table(x_test_cv$GradeClass, test_pred_rf1)
    test_error_rate_rf1=1 -sum(diag(test_conf_matrix_rf1)) / sum(test_conf_ma
trix_rf1)
    test_error_rates_rf1=c(test_error_rates_rf1, test_error_rate_rf1)

    **** model2
    rf2 = randomForest(as.factor(GradeClass) ~ ., data = x_train_cv,
                      mtry = sqrt(ncol(x_train_cv)-1), ntree = nseq[p])

    #train
    train_pred_rf2=predict(rf2, newdata = x_train_cv)
```

```

train_conf_matrix_rf2=table(x_train_cv$GradeClass, train_pred_rf2)
train_error_rate_rf2=1 -sum(diag(train_conf_matrix_rf2)) / sum(train_conf
_matrix_rf2)
train_error_rates_rf2=c(train_error_rates_rf2, train_error_rate_rf2)
#test
test_pred_rf2=predict(rf2, newdata = x_test_cv)
test_conf_matrix_rf2=table(x_test_cv$GradeClass, test_pred_rf2)
test_error_rate_rf2=1 -sum(diag(test_conf_matrix_rf2)) / sum(test_conf_ma
trix_rf2)
test_error_rates_rf2=c(test_error_rates_rf2, test_error_rate_rf2)

**** model3
rf3 = randomForest(as.factor(GradeClass) ~ ., data = x_train_cv,
                   mtry = sqrt(ncol(x_train_cv)-1)+2, ntree = nseq[p])
#train
train_pred_rf3=predict(rf3, newdata = x_train_cv)
train_conf_matrix_rf3=table(x_train_cv$GradeClass, train_pred_rf3)
train_error_rate_rf3=1 -sum(diag(train_conf_matrix_rf3)) / sum(train_conf
_matrix_rf3)
train_error_rates_rf3=c(train_error_rates_rf3, train_error_rate_rf3)
#test
test_pred_rf3=predict(rf3, newdata = x_test_cv)
test_conf_matrix_rf3=table(x_test_cv$GradeClass, test_pred_rf3)
test_error_rate_rf3=1 -sum(diag(test_conf_matrix_rf3)) / sum(test_conf_ma
trix_rf3)
test_error_rates_rf3=c(test_error_rates_rf3, test_error_rate_rf3)

**** model4
rf4 = randomForest(as.factor(GradeClass) ~ ., data = x_train_cv,
                   mtry = sqrt(ncol(x_train_cv)-1)-2, ntree = nseq[p])
#train
train_pred_rf4=predict(rf4, newdata = x_train_cv)
train_conf_matrix_rf4=table(x_train_cv$GradeClass, train_pred_rf4)
train_error_rate_rf4=1 -sum(diag(train_conf_matrix_rf4)) / sum(train_conf
_matrix_rf4)
train_error_rates_rf4=c(train_error_rates_rf4, train_error_rate_rf4)
#test
test_pred_rf4=predict(rf4, newdata = x_test_cv)
test_conf_matrix_rf4=table(x_test_cv$GradeClass, test_pred_rf4)
test_error_rate_rf4=1 -sum(diag(test_conf_matrix_rf4)) / sum(test_conf_ma
trix_rf4)
test_error_rates_rf4=c(test_error_rates_rf4, test_error_rate_rf4)

}
Mean_train_Error_Rate_rf1[p]= mean(train_error_rates_rf1)
Mean_train_Error_Rate_rf2[p]= mean(train_error_rates_rf2)

```

```

Mean_train_Error_Rate_rf3[p]= mean(train_error_rates_rf3)
Mean_train_Error_Rate_rf4[p]= mean(train_error_rates_rf4)

Mean_test_Error_Rate_rf1[p]= mean(test_error_rates_rf1)
Mean_test_Error_Rate_rf2[p]= mean(test_error_rates_rf2)
Mean_test_Error_Rate_rf3[p]= mean(test_error_rates_rf3)
Mean_test_Error_Rate_rf4[p]= mean(test_error_rates_rf4)

### full data models for oob
rf1_oob=randomForest(as.factor(GradeClass) ~ ., data = mydata,
                      mtry =(ncol(x_train_cv)-1)/2, ntree = nseq[p])
oob_rf1[p]=rf1_oob$err.rate[rf1_oob$ntree,1]

rf2_oob=randomForest(as.factor(GradeClass) ~ ., data = mydata,
                      mtry =sqrt(ncol(x_train_cv)-1), ntree = nseq[p])
oob_rf2[p]=rf2_oob$err.rate[rf2_oob$ntree,1]

rf3_oob=randomForest(as.factor(GradeClass) ~ ., data = mydata,
                      mtry =sqrt(ncol(x_train_cv)-1)+2, ntree = nseq[p])
oob_rf3[p]=rf3_oob$err.rate[rf3_oob$ntree,1]

rf4_oob=randomForest(as.factor(GradeClass) ~ ., data = mydata,
                      mtry =sqrt(ncol(x_train_cv)-1)-2, ntree = nseq[p])
oob_rf4[p]=rf4_oob$err.rate[rf4_oob$ntree,1]

}

```

Plots of error rates of random forests with different mtry values.

```

plot(nseq,Mean_train_Error_Rate_rf1,type="l",col="red",ylim=c(0,1),lwd=1.5,ma
in="Error vs number of trees for m=p/2 predictors",xlab="Number of trees",yla
b="Error")
lines(nseq,Mean_test_Error_Rate_rf1,type="l",col="black",lwd=1.5)
lines(nseq,oob_rf1,type="l",col="green",lwd=1.5)
legend("topright",c("training error","testing error","OOB error"),col=c("red"
,"black","green"),lty=c(1,1,1))

```

## Error vs number of trees for $m=p/2$ predictors



```
plot(nseq, Mean_train_Error_Rate_rf2, type="l", col="red", ylim=c(0,1), lwd=1.5, main="Error vs number of trees for m=sqrt(p) predictors", xlab="Number of trees", ylab="Error")
lines(nseq, Mean_test_Error_Rate_rf2, type="l", col="black", lwd=1.5)
lines(nseq, oob_rf2, type="l", col="green", lwd=1.5)
legend("topright", c("training error", "testing error", "OOB error"), col=c("red", "black", "green"), lty=c(1,1,1))
```

## Error vs number of trees for $m=\sqrt{p}$ predictors



```
plot(nseq, Mean_train_Error_Rate_rf3, type="l", col="red", ylim=c(0,1), lwd=1.5, main="Error vs number of trees for m=sqrt(p)+2 predictors", xlab="Number of trees", ylab="Error")
lines(nseq, Mean_test_Error_Rate_rf3, type="l", col="black", lwd=1.5)
lines(nseq, oob_rf3, type="l", col="green", lwd=1.5)
legend("topright", c("training error", "testing error", "OOB error"), col=c("red", "black", "green"), lty=c(1,1,1))
```


## Error vs number of trees for $m=\sqrt{p}+2$ predictors



```
plot(nseq, Mean_train_Error_Rate_rf4, type="l", col="red", ylim=c(0,1), lwd=1.5, main="Error vs number of trees for m=sqrt(p)-2 predictors", xlab="Number of trees", ylab="Error")
lines(nseq, Mean_test_Error_Rate_rf4, type="l", col="black", lwd=1.5)
lines(nseq, oob_rf4, type="l", col="green", lwd=1.5)

legend("topright", c("training error", "testing error", "OOB error"), col=c("red", "black", "green"), lty=c(1,1,1))
```


## Error vs number of trees for $m=\sqrt{p}-2$ predictor




Plot for test error rate of various random forest models.

```
plot(nseq, Mean_test_Error_Rate_rf1, type="l", col="blue", lwd=1.5, ylim=c(0.25, 0.55),
     main="rf with diff m", xlab="Number of trees", ylab=" Test Error Rate")
lines(nseq, Mean_test_Error_Rate_rf2, type="l", col="black", lwd=1.5)
lines(nseq, Mean_test_Error_Rate_rf3, type="l", col="green", lwd=1.5)
lines(nseq, Mean_test_Error_Rate_rf4, type="l", col="orange", lwd=1.5)
legend("topright", c("m=p/2", "m=sqrt(p)", "m=sqrt(p)+2", "m=sqrt(p)-2"), col=c("blue", "black", "green", "orange"), lty=rep(1, 4))
```





```
plot(nseq, Mean_test_Error_Rate_rf1, type="l", col="blue", lwd=1.5, ylim=c(0.27, 0.3),
     main="rf with diff m", xlab="Number of trees", ylab="Test Error Rate")
lines(nseq, Mean_test_Error_Rate_rf3, type="l", col="green", lwd=1.5)
```



To get optimal number of trees

```
tune_rf=tune.randomForest(as.factor(GradeClass)~.,data=train,
                           ntree=seq(10,200,by=10),
                           mtry=c((ncol(train)-1)/2,sqrt(ncol(train)-1),
                                   sqrt(ncol(train)-1)+2,sqrt(ncol(train)-1)-2))
summary(tune_rf)

##
## Parameter tuning of 'randomForest':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   mtry ntree
##     6   190
##
## - best performance: 0.2733961
##
## - Detailed performance results:
##       mtry ntree   error dispersion
## 1  6.000000    10 0.2959745 0.02625807
## 2  3.464102    10 0.3268985 0.04120148
## 3  5.464102    10 0.3089418 0.02294025
## 4  1.464102    10 0.4757165 0.04288541
## 5  6.000000    20 0.3001447 0.02024327
## 6  3.464102    20 0.2955439 0.03206686
## 7  5.464102    20 0.2934467 0.04041309
## 8  1.464102    20 0.4920171 0.03655764
## 9  6.000000    30 0.2880300 0.02492813
## 10 3.464102    30 0.2922071 0.02858565
## 11 5.464102    30 0.2926151 0.04133638
## 12 1.464102    30 0.4920153 0.03959334
## 13 6.000000    40 0.2834153 0.03613240
## 14 3.464102    40 0.2901046 0.03040444
## 15 5.464102    40 0.2825785 0.03458402
## 16 1.464102    40 0.4936890 0.03720399
## 17 6.000000    50 0.2817381 0.03141690
## 18 3.464102    50 0.2947089 0.02412825
## 19 5.464102    50 0.2792503 0.03129301
## 20 1.464102    50 0.4928539 0.03669745
## 21 6.000000    60 0.2754759 0.03164478
## 22 3.464102    60 0.2901168 0.02549357
## 23 5.464102    60 0.2809205 0.02509367
## 24 1.464102    60 0.4932706 0.03712242
## 25 6.000000    70 0.2805056 0.02731655
## 26 3.464102    70 0.2989017 0.03306396
## 27 5.464102    70 0.2788302 0.03016615
## 28 1.464102    70 0.4936890 0.03720399
## 29 6.000000    80 0.2779934 0.02966311
```

## 30	3.464102	80	0.2876046	0.03665327
## 31	5.464102	80	0.2784083	0.02985490
## 32	1.464102	80	0.4932706	0.03790026
## 33	6.000000	90	0.2758996	0.03537375
## 34	3.464102	90	0.2917870	0.02719662
## 35	5.464102	90	0.2800819	0.02728760
## 36	1.464102	90	0.4936890	0.03720399
## 37	6.000000	100	0.2742225	0.03095104
## 38	3.464102	100	0.2942992	0.02692673
## 39	5.464102	100	0.2842591	0.02930796
## 40	1.464102	100	0.4936890	0.03720399
## 41	6.000000	110	0.2767277	0.02927252
## 42	3.464102	110	0.2846810	0.03161181
## 43	5.464102	110	0.2759066	0.03283934
## 44	1.464102	110	0.4936890	0.03720399
## 45	6.000000	120	0.2767347	0.02923655
## 46	3.464102	120	0.2880160	0.02699396
## 47	5.464102	120	0.2775732	0.03291637
## 48	1.464102	120	0.4936890	0.03720399
## 49	6.000000	130	0.2742225	0.03650623
## 50	3.464102	130	0.2951360	0.03055609
## 51	5.464102	130	0.2738075	0.03108318
## 52	1.464102	130	0.4936890	0.03720399
## 53	6.000000	140	0.2746496	0.02836638
## 54	3.464102	140	0.2855265	0.02285123
## 55	5.464102	140	0.2746409	0.03219766
## 56	1.464102	140	0.4936890	0.03720399
## 57	6.000000	150	0.2821653	0.03005785
## 58	3.464102	150	0.2813476	0.02772553
## 59	5.464102	150	0.2784170	0.02802461
## 60	1.464102	150	0.4936890	0.03720399
## 61	6.000000	160	0.2750662	0.02679088
## 62	3.464102	160	0.2821775	0.02501423
## 63	5.464102	160	0.2738128	0.02580860
## 64	1.464102	160	0.4936890	0.03720399
## 65	6.000000	170	0.2742242	0.02717717
## 66	3.464102	170	0.2813389	0.02489467
## 67	5.464102	170	0.2742347	0.02629716
## 68	1.464102	170	0.4936890	0.03720399
## 69	6.000000	180	0.2763145	0.02624992
## 70	3.464102	180	0.2905370	0.02671991
## 71	5.464102	180	0.2788337	0.02461305
## 72	1.464102	180	0.4936890	0.03720399
## 73	<b>6.000000</b>	<b>190</b>	<b>0.2733961</b>	<b>0.02319444</b>
## 74	3.464102	190	0.2855143	0.03017368
## 75	5.464102	190	0.2800837	0.02995617
## 76	1.464102	190	0.4936890	0.03720399
## 77	6.000000	200	0.2763128	0.02688728
## 78	3.464102	200	0.2834222	0.02873554

```
## 79 5.464102    200 0.2754881 0.02738155
## 80 1.464102    200 0.4936890 0.03720399
```

Fitting model based on tuned parameters.

```
rf_final_train=randomForest(as.factor(GradeClass)~.,data=train,
                             mtry=6,ntree=190)
print(rf_final_train)


##
## Call:
## randomForest(formula = as.factor(GradeClass) ~ ., data = train,      mtry
## = 6, ntree = 190)
##              Type of random forest: classification
##              Number of trees: 190
##
## No. of variables tried at each split: 6
##
##              OOB estimate of  error rate: 27.96%
## Confusion matrix:
##      A  B   C   D   F class.error
## A 20 32   6   8  12   0.7435897
## B 16 85  52   7  15   0.5142857
## C  0 34 157  66  16   0.4249084
## D  1  8  58 137  76   0.5107143
## F  1  5   5  50 807   0.0702765

test_pred_rf_final=predict(rf_final_train, newdata =test)
test_conf_matrix_rf_final=table(test$GradeClass, test_pred_rf_final)
test_error_rate_rf_final=1 -sum(diag(test_conf_matrix_rf_final)) / sum(test_c
onf_matrix_rf_final)
test_error_rate_rf_final

## [1] 0.2952646

barplot(sort(importance(rf_final)[,1]),horiz=T,
        las=1,cex.names = 0.4,main="Variable Importance Plot")
```

## Variable Importance Plot



Method	m	Best n tree	OOB error	Testing error
Random forest	p/2=6	150	0.2796	0.2952

## 7.Boosting

```
library(gbm)

boost=gbm(as.factor(GradeClass) ~., data =train,
          distribution = "multinomial", n.trees = 500,
          interaction.depth = 2, shrinkage = 0.01)


train_pred_class=apply(predict.gbm(boost,train,type="response"),1,which.max)-1

## Using 500 trees...

test_pred_class=apply(predict.gbm(boost,test,type="response"),1,which.max)-1

## Using 500 trees...

barplot(sort(importance(rf_final)[,1]),horiz=T,
        las=1,cex.names = 0.4,main="Variable Importance Plot")
```



Train and Test error rate.

```
train_conf_mat=table(train$GradeClass,train_pred_class)
test_conf_mat=table(test$GradeClass,test_pred_class)

test_accuracy=sum(diag(test_conf_mat))/sum(test_conf_mat)
```

```
test_err_rate_boosting=1-test_accuracy
paste("test error:",test_err_rate_boosting)

## [1] "test error: 0.309192200557103"

train_accuracy=sum(diag(train_conf_mat))/sum(train_conf_mat)
train_err_rate_boosting=1-train_accuracy
paste("train error:",train_err_rate_boosting)

## [1] "train error: 0.241935483870968"
```

## Conclusion:

Method	Testing Accuracy
Multinomial Logistic Regression	0.7325905
KNN	0.6866295
SVM	0.7311977
Decision tree	0.6824512
Bagging	0.7019499
Random Forest	0.7047354
Boosting	0.6908078